Pivotum

THE DEGREE RISK INDEX · FALL 2026 EDITION

Will AI Replace Teachers?

The complete profile — the full scores by track, the six factors behind them, the honest downsides, and what to ask before you commit.

3.6

AI EXPOSURE · ELEMENTARY / EARLY YEARS

where 10 is most at risk

Read this one first; send the student version to them.

HIGHLIGHTING IS COLOUR-CODED

THE FINDING EXPOSURE PROTECTION METHOD & HONESTY

Before you read this — a note for parents

You're likely reading this before your child is. That's normal, and it's worth pausing on how you use it.

Don't lead with the scores. If you open with a risk number, they either get frightened and shut down, or defensive about a decision they've already made.

What this is. A facts document, not a recommendation.

A practical note: there's a separate short version written directly to them. Send them that one.

And one thing specific to teaching. This is the profile where the AI answer and the career answer point in opposite directions, and it would be easy to conflate them. AI is genuinely good news for classroom teachers — it removes the parts of the job people leave over. The reasons teaching is hard have nothing to do with technology and haven't changed.

If your instinct on hearing "teaching" was to worry about pay and workload rather than robots, that instinct is correct. This document will confirm it, and it will also show you something more useful: that the difficulty varies more by *where* someone teaches than by whether they teach.

The short answer

Yes — and AI is, on balance, actively good news for classroom teachers.

Classroom teaching scores **3.6** out of 10. Special education scores **3.1**. But online instruction and content-based tutoring score **6.8**, nearly double.

The thing that makes teaching interesting analytically is this: teaching rates higher on automatability than nursing or medicine — and still lands in the low band. Understanding why is the most useful thing in this profile, and it explains what actually protects work.

The scores

SETTING	2023	2025	FALL 2026	3-YR MOVEMENT	BAND
Special education	2.7	2.9	**3.1**	+0.4	Low
Elementary / early years	3.2	3.4	**3.6**	+0.4	Low
Secondary — STEM subjects	3.4	3.6	**3.8**	+0.4	Low
Secondary — general subjects	3.7	3.9	**4.1**	+0.4	Low–Moderate
Online instruction / tutoring	5.5	6.2	**6.8**	+1.3	Moderate–High

Scores run 1–10, where 10 is most at risk. For context: the median profession in this edition sits around 5.5, bedside nursing scores 2.8, licensed engineering 4.0, and entry-level software development 8.1.

Read the table this way: classroom teaching's +0.4 over three years is tied for the slowest movement in the entire index. Online instruction moved +1.3 over the same period — the same subject knowledge, the same automatable tasks, on a completely different trajectory.

Why teaching scores low — the six factors

How much of this job can AI already do?	**7.7**
How hard will it be to land that first job?	**1.5**
Does it have to be done in person, with your hands?	**9.0**
Does someone need a human they can trust and hold responsible?	**9.5**
Does the law require a licensed human?	**8.0**
How often does the job hit genuinely new, high-stakes situations?	**7.5**

EXPOSURE RAISES THE SCORE   PROTECTION LOWERS IT

Here's how elementary classroom teaching answers them.

HOW MUCH OF THIS JOB CAN AI ALREADY DO? (35% OF THE SCORE)

A great deal — more than nursing, more than medicine, more than most careers we score.

Lesson planning and resource creation. Marking and feedback generation. Differentiating materials by level. Parent communication drafts. Administrative reporting. Curriculum mapping.

We rate this **7.7**. Section 5 explains why a rating that high still produces a score of 3.6, because it's the clearest demonstration of how this whole framework works.

HOW HARD WILL IT BE TO LAND THAT FIRST JOB? (15%)

The strongest protection we give on this factor anywhere.

Certification is mandated and requires supervised classroom practice. Shortages are widespread — NCES reports that 74% of public schools had difficulty filling at least one vacancy, and shortages are acute in mathematics, science, special education and bilingual instruction.

A newly qualified teacher in a shortage subject enters one of the more welcoming graduate markets in this index.

DOES IT HAVE TO BE DONE IN PERSON, WITH YOUR HANDS? (15%)

Substantially — and this is where the internal split comes from.

A classroom is embodied, social work. Thirty children in a room is a physical management problem before it is an instructional one.

DOES SOMEONE NEED A HUMAN THEY CAN TRUST AND HOLD RESPONSIBLE? (15%)

Yes, and in an unusually deep form.

Parents hand over their children. That is not a relationship anyone delegates to software, and safeguarding duties attach to a named adult.

DOES THE LAW REQUIRE A LICENSED HUMAN? (10%)

Yes — certification is mandated, and supervision of minors is legally reserved.

HOW OFTEN DOES THE JOB HIT GENUINELY NEW, HIGH-STAKES SITUATIONS? (10%)

Frequently. Every class is a different mix of children on a different day, and the judgment calls — behavioural, pastoral, safeguarding — arrive without warning.

Why a 7.7 automatability rating still produces a 3.6

This is the most instructive section in the profile, and the argument generalises well beyond teaching.

Automatability carries the most weight in our model at 35% — more than any other single factor. Teaching rates **7.7** on it, higher than nursing, higher than medicine, higher than most careers we score. A very large share of a teacher's working hours goes to planning, resource creation, differentiation, marking, feedback and reporting. Nearly all of it is now automatable, and much is already being automated in practice.

If our model let automatability alone drive the score, teaching would sit in the moderate-to-high band alongside marketing and business analysis.

It doesn't, because the other five factors run strongly the other way.

Presence rates 9.0 — a classroom is embodied, social work. Trust rates 9.5, because parents hand over their children. Certification is mandated and requires classroom practice. And schools perform a custodial and social function that has nothing to do with information transfer at all.

The design choice producing this result is deliberate. Our three protection factors carry 50% of the weight in total — more than automatability's 35%. That means protection can outweigh exposure. If we had weighted exposure more heavily, teaching would score high and the model would be wrong.

And it is exactly why the online track scores 6.8. Take the identical subject knowledge and the identical automatable tasks. Strip out physical presence. Weaken trust. Remove the certification requirement. Now the same 7.7 automatability rating dominates, because there is nothing left to counterweight it.

Nothing about the *work* changed. The setting did. That single comparison is the best evidence in this index that where you do something matters more than what you know.

What the trend shows

Classroom teaching: 3.2 → 3.4 → 3.6. Up 0.4 in three years, and every point of it is administrative automation.

Online instruction: 5.5 → 6.2 → 6.8. Up 1.3.

The gap widened from 2.3 points to 3.2 — one of the fastest-widening splits we measure, and it is entirely about delivery mode rather than subject or skill.

How durable is the protection?

PROTECTION	STRENGTH TODAY	DURABILITY	WHERE THE PRESSURE IS COMING FROM
Human trust / accountability	**Among the highest we score**	**Most durable**	Parents handing over children is socially rooted and legally reinforced by safeguarding duty.
Regulatory / licensing	Strong	**Durable**	Certification and supervision-of-minors requirements move at the speed of legislation.
Physical / embodied	Strong	**Durable in a classroom**	A room of thirty children is a management problem before an instructional one.
Judgment under novelty	Good	**Eroding slowly**	Pastoral and behavioural judgment is not automating; instructional planning is.

The honest read. Teaching's protections are unusually well matched to the threat. The thing AI does best — generating instructional content — is the thing schools need least help with. The things schools actually need, a trusted adult managing a room of children, are where AI is weakest.

How the role will actually change

WAS THE JOB	IS BECOMING THE JOB
Planning lessons from scratch	Adapting generated plans to this class
Marking a set of thirty	Reviewing generated feedback and adding what matters
Differentiating by hand for three ability groups	Differentiation arrives automatically
Writing reports over half-term	Drafted, then checked and personalised
Managing thirty children in a room	**Unchanged**
Noticing the child who's gone quiet	**Unchanged**

The pattern across this index is that AI absorbs the routine, repeatable layer first. In creative and knowledge work, that layer was often the craft itself. In teaching it is the evening and weekend work — which is the single most-cited reason people leave.

Two honest caveats. Verification becomes a real professional responsibility — generated feedback that misreads a child's work is worse than no feedback. And the time saved may be absorbed by larger classes or additional duties rather than returned to teachers, which is a funding decision rather than a technological one.

Human strengths that will matter most

1. **Managing a room** — thirty children, social dynamics, attention, behaviour. The core of

> the job and the least automatable thing in it.

2. **Noticing** — the child who has gone quiet this week, whose home situation has changed.

> Pastoral judgment that no system has visibility of.

3. **Being trusted** — by children and by parents. Teaching's deepest protection. 4. **Knowing why this student is stuck** — diagnosis, not delivery. 5. **Verification** — knowing when generated

feedback or a generated plan is wrong for this

> class.

Notably absent: subject knowledge recall and resource creation. These were traditional markers of a strong early-career teacher and are the fastest-depreciating.

Other things to consider about this job

What's genuinely good about it

Adoption of AI in classrooms is already near-universal, and — unusually for this index — teachers largely report it as a relief. The tasks being automated are the ones people cite when they leave.

What teachers report as rewarding is consistent:

- **Watching a child understand something** — the thing nearly everyone names first
- **Relationships that build over a year** and sometimes over a career
- **Autonomy inside the classroom**, which is higher than most jobs at equivalent pay
- **Genuine social contribution**, which people mean rather than say
- **Rhythm and predictability** — terms, holidays, a schedule that aligns with family life

The practical case: certification is portable, shortages mean employability is high especially in STEM and special education, and the pension provision in most public systems is better than the private-sector equivalent.

What's genuinely hard about it

The workload is the headline complaint, and the visible hours are a fraction of the real ones. Planning, marking and reporting have historically consumed evenings and weekends — which is precisely why the automation in section 7 matters.

Pay lags comparable graduate careers, particularly mid-career, and that gap is the most-cited reason for leaving.

Behaviour management is harder than most students expect, and it is the thing teacher training prepares people for least well.

And the attrition is real — though the direction of travel is better than the coverage suggests. RAND's tracking of district-reported turnover, now in its fifth year, finds educator turnover continuing to decline toward pre-pandemic levels. That is the opposite of the popular narrative and worth knowing.

Set against that, the Learning Policy Institute's analysis of national data finds turnover has been persistently high for two decades and exceeds that of other high-performing education systems. Both are true: it is high by international standards and currently improving.

Who this work suits — and who it doesn't

Teaching tends to suit people who:

- **Can hold a room.** The single most important trait and the hardest to teach.
- **Get energy from other people** rather than being drained by them — thirty of them, all day
- **Are patient with slow progress**, and can explain the same thing eight different ways
- **Are organised**, because the administrative load is real even after automation
- **Want their work to matter visibly**

It's harder for people who:

- Need **quiet to function**
- Find **conflict draining** — behaviour management is unavoidable
- Want **high autonomy over their time**; the school day is fixed
- Are drawn primarily to their **subject** rather than to children. This is the common mistake, and it's a genuinely different motivation.

What else is moving this market

Shortages are structural, and worst in mathematics, science, special education and bilingual instruction. Choosing a shortage subject materially improves employability, pay and geographic freedom.

Where you teach matters more than almost anything. The Learning Policy Institute's analysis of national data found turnover of 19.1% in high-poverty schools against 12.9% in low-poverty schools. That is a larger difference than between many entirely separate professions.

Working conditions move with funding and politics rather than technology. A student should understand that the *technology* outlook for teaching is genuinely positive and the *conditions* outlook depends on decisions nobody in this document controls.

The AI-native advantage — what to actually do about it

The situation: classroom AI adoption is near-universal, and it is arriving faster than the training to use it well. Under half of teachers report receiving any institutional guidance.

That gap is the opportunity. A newly qualified teacher who is genuinely fluent walks into a shortage market with something most experienced colleagues don't have.

WHAT AN AI-NATIVE TEACHER LOOKS LIKE

1. Verify. Knowing when generated feedback misreads a child's work, or when a generated plan doesn't fit this class. This is a pedagogical skill rather than a technical one — you can only catch a plausible error if you know the children.

2. Reclaim the time deliberately. The automation only helps if the hours saved go somewhere. Teachers who are intentional about this get the benefit; those who aren't watch it fill with something else.

3. Teach students to use it honestly. Every child in the room is already using these tools. A teacher who can set out what's cheating and what's learning — rather than banning and hoping — is doing something schools urgently need.

4. Assess for understanding rather than output. If a generated essay can pass your assessment, your assessment was measuring the wrong thing. This is the biggest unsolved problem in school assessment and the profession needs practitioners who can think about it.

5. Contribute to policy. Schools are writing AI rules right now, mostly without practitioner input.

CONCRETE PREPARATION

Before training: ask providers how AI appears in the curriculum. The answers vary enormously and are revealing.

During: get classroom exposure early and often — behaviour management is the thing people underestimate. Practise the verification habit deliberately.

Choosing a subject: shortage subjects — mathematics, physics, chemistry, special education, bilingual — carry better employability, better pay and more geographic freedom. That decision matters more than most trainees realise.

The routes in

ROUTE	TYPICAL LENGTH	NOTES
Education degree	3–4 years	Direct route from school.
Subject degree + teaching qualification	3–4 years + 1	Common, and strong in shortage subjects.
Employment-based / school-direct training	1–2 years	Earn while training. Expanding, and worth investigating.
Career-change routes	1–2 years	Widely available and often subsidised in shortage subjects.
Teaching assistant → qualified	Varies	Slower, but arrives with real classroom experience.

Where a teaching qualification can lead later

Broad, and more portable than people assume: classroom practice across phases, special education, subject leadership and headship, curriculum and assessment design, educational technology, teacher training, policy, and international schools.

One caution, and it's the same pattern as everywhere in this index. Education-adjacent careers built on content delivery — tutoring businesses, online course instruction, test prep — share the online track's exposure at 6.8 rather than the classroom's protection at 3.6. They look like teaching and they score like content production.

What to look for in a teacher training program

1. Time in classrooms, and how early. The single best predictor of whether someone stays. Ask how many weeks, across how many different schools.

2. Behaviour management taught properly. Not a module — sustained, practical, with real feedback. It's the thing new teachers report being least prepared for.

3. Is AI treated as part of practice? Students should learn to use these tools *and* to check them, and to think about assessment design in a world where generated work is free.

4. Placement variety. Different phases, different intakes, different school contexts. Given the 19.1%-versus-12.9% turnover gap, understanding what different settings are actually like matters enormously.

5. Retention data. What proportion of their graduates are still teaching at three and five years?

Questions to ask a teacher training provider

On classroom time:

- How many weeks in schools, in which years, and across how many settings?
- Do trainees teach in both high- and low-need schools?

On preparation:

- How is behaviour management taught, and how much practical feedback do trainees get?
- How does the programme prepare trainees to work with AI tools, and to redesign assessment?

On outcomes:

- What proportion of your graduates are still teaching at three years? At five?
- Which subjects do you place best in?

Red flags: minimal or late classroom exposure; behaviour management as a single module; no retention data; a curriculum unchanged since 2022; no engagement with AI beyond a policy banning it.

Go deeper — further reading

- **Teacher Turnover in the United States: Patterns, Drivers, and Policy Strategies** — *Learning Policy Institute*, drawing on NCES national surveys. The primary source for the turnover figures, including the high-poverty versus low-poverty gap. The most rigorous treatment available and free.

- **Educator Turnover Continues Decline Toward Prepandemic Levels** — *RAND Corporation*, *American School District Panel*. Five years of district-reported tracking showing turnover falling. Read it alongside the LPI factsheet: they measure different things and both are right.

- **What's It Like To Be A Teacher in America Today?** — *Pew Research Center*. Teacher-reported satisfaction, workload and intent to leave, from a research organisation rather than an advocacy body.

The counterargument — read this one:

> The strongest case against our 3.6 isn't that AI will replace classroom teachers. It's that we may be measuring the wrong risk entirely.
>
> On this reading, teaching's problem was never automation — it was pay, workload, behaviour and status. A framework that scores AI exposure will always make teaching look safe while saying nothing about the thing that actually drives people out at 5.3% vacancy rates and double-digit turnover. A parent reading "3.6, low risk" could reasonably conclude teaching is a safe bet, when the profession's real attrition problem is untouched by anything in our six factors.
>
> **Where we land:** this is entirely fair, and it's why section 9 exists. Our score measures automation exposure and nothing else. For teaching specifically, the AI answer is genuinely good news and the career answer is genuinely complicated, and conflating them would be the most misleading thing this document could do. We'd rather a reader took the 3.6 as "technology is not the risk here" than as "this is easy."

Bottom line

Teaching is well protected, and AI is a genuine improvement to the job rather than a threat to it.

The protections are unusually well matched to the threat: the thing AI does best is generate instructional content, which is the thing schools need least help with. What schools actually need — a trusted adult managing a room of children, noticing what nobody reported — is where AI is weakest.

The most useful finding here is analytical rather than reassuring. Teaching rates 7.7 on automatability and still scores 3.6, which shows what actually protects work: not the absence of automatable tasks, but the presence of everything else. Take the same knowledge online and the score nearly doubles.

The real risks are the ones that were always there. Workload, pay, behaviour, and the fact that turnover in a high-poverty school runs at 19.1% against 12.9% in a low-poverty one. AI helps meaningfully with the first of those and not at all with the rest.

So the useful advice is: choose a shortage subject, interrogate a training provider's classroom hours and retention data rather than its ranking, be deliberate about reclaiming the time automation gives back, and understand that *where* you teach will shape your experience more than *whether* you teach.

And it's worth stating what the score can't measure. Watching a child understand something they couldn't do last week is a thing very few jobs offer at all. The difficulties above are real. So is that.

Discussion questions

Part 1 — About the work itself

1. What made teaching appealing? Push past "I like my subject" or "I want to help" to something specific. 2. **Are they drawn to children, or to their subject?** This is the most important question here, and the two lead to very different experiences. 3. Can they hold a room? Have they ever had to — coaching, leading, running something? 4. What does a Tuesday look like in the job they're imagining?

Part 2 — Reacting to the findings

5. Classroom teaching scores 3.6 and online instruction 6.8. Was that surprising? 6. The profile says teaching rates *higher* on automatability than nursing or medicine and still scores low. Does

that make sense to them? 7. AI takes planning, marking and differentiation. Does that change how they feel about the job?

Part 2b — The harder conversation

8. **Pay.** Have you looked at actual teacher salaries at year one, year five and year fifteen where you live, against other graduate careers? 9. Behaviour management is the thing new teachers report being least prepared for. What do they think their instinct would be? 10. Turnover runs at 19.1% in high-poverty schools against 12.9% in low-poverty ones. Does that change how they'd think about where to work?

Part 2c — The other side

11. Of what teachers say is rewarding — watching a child understand, relationships over a year, classroom autonomy, the rhythm of the year — which two matter most? 12. Looking at "who this work suits": which items sound like them? Answer separately, then compare. 13. Do thirty people all day energise them or drain them? Honest answer.

Part 2d — The AI question, practically

14. Every child in their future classroom will be using these tools. How would they want to handle that? 15. The profile says if a generated essay can pass your assessment, the assessment was measuring the wrong thing. What do they make of that?

Part 3 — Testing it against reality

16. **The most valuable single action:** get into a classroom. Volunteering, assisting, anything. Nothing else answers the fit question as fast. 17. Find a teacher three to five years in — not one about to retire — and ask what a normal week actually looks like.

Part 4 — About the program

18. Ask each provider the classroom-hours and retention questions from section 14. 19. How is behaviour management actually taught?

Part 5 — Widening the frame

20. **Which subject?** Shortage subjects — maths, physics, chemistry, special education, bilingual — carry better employability, pay and geographic freedom. 21. Have they considered special education? It scores best in the profile at 3.1 and has the deepest shortages. 22. If teaching

weren't available, what would they choose? The answer usually reveals what they're actually drawn to.

Part 6 — For the parent, alone

23. What was my honest reaction when they said teaching? Where did that come from? 24. Am I worried about the money, and have I said so clearly rather than hinting? 25. What's the one thing I most want them to understand — in a sentence, without a number in it?

This is analysis and informed judgment about a fast-moving situation. For teaching specifically, we've flagged that our score measures automation exposure only — and that the profession's real attrition drivers sit entirely outside what our six factors can see.

Scores for 2023 and 2025 are retrospectively calculated using the current methodology. Scores measure exposure to what AI can already do — not how much any particular employer has deployed.

Technical scoring appendix — Teaching

Pivotum Degree Risk Index — Fall 2026 Edition

This appendix is the audit trail. It sets out the formula, the rating anchors, every factor rating behind every track score in this profile, the backcast arithmetic, a sensitivity analysis, and the specific things that would change our mind.

We publish it because a score nobody can check is an opinion with a decimal point.

1. The formula

```
Risk = 0.35 × Automatability
     + 0.15 × EntryErosion
     + 0.15 × (10 − Physical)
     + 0.15 × (10 − Trust)
     + 0.10 × (10 − Regulatory)
     + 0.10 × (10 − Judgment)
```

Two exposure factors carry 50% of the weight; three protection factors carry 40%; judgment under novelty carries the remaining 10% and sits on the protection side. The deliberate consequence is that protection can outweigh exposure — which is why a highly automatable job like teaching still scores low.

Bands. Very low below 3.0 · Low 3.0–4.4 · Moderate 4.5–5.4 · Moderate–High 5.5–6.9 · High 7.0 and above.

2. Rating anchors

HOW THE 0–10 RAW RATINGS ARE ANCHORED

Every factor is rated on its own 0–10 scale before weighting. The anchors are identical across all twenty-seven careers, which is what makes the scores comparable.

Task automatability (A) — *what share of the working week could a capable current system already do to an acceptable standard?*

RATING	ANCHOR
0–2	Almost nothing. The work is physical, relational or wholly novel.
3–4	Administrative surround only — notes, scheduling, correspondence.
5–6	A meaningful minority of the week, mostly documentation and lookup.
7–8	A large share, including tasks central to the junior version of the role.
9–10	Most of the described duties, at or near professional standard.

Entry-path erosion (E) — *how much has the traditional route into this work narrowed?*

RATING	ANCHOR
0–2	Protected by statute or physical necessity. Training hours cannot be automated.
3–4	Stable. Some junior tasks absorbed; the apprenticeship survives.
5–6	Visible narrowing. Employers prefer experience; junior intake reduced.
7–8	Substantial. The work juniors learned on has largely gone.
9–10	The training layer and the automatable layer were the same layer.

Physical / embodied (P) — *must this be done in person, with a body, in an unpredictable setting?*

RATING	ANCHOR
0–2	Fully screen-based. Location-independent.
3–4	Occasional presence; the core work is not physical.
5–6	Regular physical presence, in largely controlled conditions.
7–8	Substantially physical, with some variability of setting.
9–10	Irreducibly physical, in conditions that differ every time.

Human trust / accountability (T) — *does someone need a specific human they can rely on, and must a human answer when it goes wrong?*

RATING	ANCHOR
0–2	Output is anonymous and reviewed by others.
3–4	Work is signed off by someone else. No personal reliance.
5–6	Some client relationship; accountability is institutional.
7–8	Named responsibility, with commercial or professional consequence.
9–10	The relationship is the service, and liability attaches personally.

Regulatory / licensing (R) — *does the law reserve this act to a qualified human*?

RATING	ANCHOR
0–2	No licence exists and none is proposed.
3–4	Certification is customary but not reserved.
5–6	Partial reservation, varying by jurisdiction or task.
7–8	Licensed practice, with some unreserved adjacent work.
9–10	Comprehensive statutory reservation, including of the training route.

Judgment under novelty (J) — *how often does the work present something with no matching precedent*, *where being wrong is costly*?

RATING	ANCHOR
0–2	Routine by design. Deviation is an error, not a case.
3–4	Occasional exceptions, handled by escalation.
5–6	Regular non-standard cases within a known range.
7–8	Frequent genuine novelty with material consequences.
9–10	Novelty is constant, and in some fields adversarially generated.

3. Factor ratings by track

Ratings are the Fall 2026 edition unless a range is shown. **P, T, R and J are held constant across editions** — those protections are structural and do not move on a three-year horizon. All measured movement is in A and E.

SPECIAL EDUCATION — 3.1 (LOW)

FACTOR	WEIGHT	2023	2025	2026	CONTRIBUTION TO 2026 SCORE
Task automatability	35%	5.7	6.2	**6.6**	2.31
Entry-path erosion	15%	1.5	1.8	**2.0**	0.3
Physical / embodied	15%	9.0	9.0	**9.0**	0.15
Human trust / accountability	15%	9.5	9.5	**9.5**	0.07
Regulatory / licensing	10%	9.0	9.0	**9.0**	0.1
Judgment under novelty	10%	8.5	8.5	**8.5**	0.15
				Total	**3.08 → 3.1**

ELEMENTARY / EARLY YEARS — 3.6 (LOW)

FACTOR	WEIGHT	2023	2025	2026	CONTRIBUTION TO 2026 SCORE
Task automatability	35%	6.9	7.4	**7.8**	2.73
Entry-path erosion	15%	1.8	2.0	**2.2**	0.33
Physical / embodied	15%	9.0	9.0	**9.0**	0.15
Human trust / accountability	15%	9.5	9.5	**9.5**	0.07
Regulatory / licensing	10%	9.0	9.0	**9.0**	0.1
Judgment under novelty	10%	8.0	8.0	**8.0**	0.2
				Total	**3.59 → 3.6**

SECONDARY — STEM SUBJECTS — 3.8 (LOW)

FACTOR	WEIGHT	2023	2025	2026	CONTRIBUTION TO 2026 SCORE
Task automatability	35%	7.2	7.7	**8.2**	2.87
Entry-path erosion	15%	1.8	2.0	**2.2**	0.33
Physical / embodied	15%	8.8	8.8	**8.8**	0.18
Human trust / accountability	15%	9.3	9.3	**9.3**	0.1
Regulatory / licensing	10%	9.0	9.0	**9.0**	0.1
Judgment under novelty	10%	7.8	7.8	**7.8**	0.22
				Total	**3.8 → 3.8**

SECONDARY — GENERAL SUBJECTS — 4.1 (LOW)

FACTOR	WEIGHT	2023	2025	2026	CONTRIBUTION TO 2026 SCORE
Task automatability	35%	7.9	8.4	**8.8**	3.08
Entry-path erosion	15%	2.0	2.2	**2.5**	0.38
Physical / embodied	15%	8.8	8.8	**8.8**	0.18
Human trust / accountability	15%	9.3	9.3	**9.3**	0.1
Regulatory / licensing	10%	9.0	9.0	**9.0**	0.1
Judgment under novelty	10%	7.5	7.5	**7.5**	0.25
				Total	**4.09 → 4.1**

ONLINE INSTRUCTION / TUTORING — 6.8 (MODERATE–HIGH)

FACTOR	WEIGHT	2023	2025	2026	CONTRIBUTION TO 2026 SCORE
Task automatability	35%	4.4	6.1	**7.6**	2.66
Entry-path erosion	15%	4.5	5.2	**5.8**	0.87
Physical / embodied	15%	1.5	1.5	**1.5**	1.27
Human trust / accountability	15%	5.0	5.0	**5.0**	0.75
Regulatory / licensing	10%	3.0	3.0	**3.0**	0.7
Judgment under novelty	10%	4.5	4.5	**4.5**	0.55
				Total	**6.8 → 6.8**

4. Backcast arithmetic

Worked example for the headline track, showing exactly how the three-year movement is composed.

Special education

EDITION	A	E	PROTECTION DEFICIT (FIXED)	SCORE
2023	5.7	1.5	0.47	**2.7**
2025	6.2	1.8	0.47	**2.9**
Fall 2026	6.6	2.0	0.47	**3.1**

Movement decomposition: automatability contributed **+0.31** and entry-path erosion **+0.07**, for a total of **+0.4** points over three years.

The 2023 and 2025 figures are retrospective reconstructions using the current methodology, not archived scores from published editions. They are our best assessment of what each factor would have rated at the time, given what was then demonstrable. They are not measurements.

5. Sensitivity

How far the score moves if a single factor is rated one point differently:

FACTOR	±1 POINT MOVES THE SCORE BY
Task automatability	±0.35
Entry-path erosion	±0.15
Physical / embodied	±0.15
Human trust / accountability	±0.15
Regulatory / licensing	±0.10
Judgment under novelty	±0.10

What this means in practice. A one-point disagreement about automatability moves a score by 0.35 — enough to shift a track between bands at the margins. A one-point disagreement about licensing moves it by 0.10, which almost never changes a band.

So if you disagree with one of our numbers, the automatability rating is where the disagreement matters most, and it is the rating we would most want challenged.

6. Stated limitations

We measure capability, not deployment. Scores reflect exposure to what AI can already do at the leading edge — not how much any particular employer has adopted. Adoption lags capability by years and lags unevenly: large, well-funded organisations first; small, rural, public-sector and thin-margin employers much later. The lag is a buffer, not a shelter. It buys time without changing direction.

We do not measure demand. A task can be highly automatable and highly in demand simultaneously. Where the labour market and our framework disagree, we say so in the profile rather than adjusting the score to fit.

We do not measure oversupply. The number of people entering a field is outside our six factors, and for some careers it matters more than automation.

We do not measure industry economics. A profession can contract for reasons entirely unrelated to technology.

We do not measure whether the work is worth doing. Pay, satisfaction, debt, hours and meaning sit outside the score and are covered separately in each profile — often at greater length, because for several careers they matter more.

Sub-track definitions are ours. Job titles vary between employers and countries. We have chosen tracks that reflect genuinely different work, not official classifications.

---

7. What would change our mind

Each edition names specific, checkable indicators. If these move, the scores move — including downward.

Across the index:

- **Graduate hiring share.** If employers in the professions where we rate entry-path erosion highly return to hiring juniors at pre-2023 rates, those ratings fall.
- **Content of mandated training hours.** Several professions protect their on-ramp through required supervised experience. If the content of those hours shifts from producing work to supervising it, the protection weakens without any rule changing.
- **Offsite and prefabricated construction share.** The clearest test of our physical-protection ratings: automation performs well in controlled settings, so moving work indoors raises exposure without any robotics breakthrough.
- **Verified deployment rather than capability.** If adoption stalls materially below capability for several editions, our leading-indicator approach is overstating near-term risk and we will say so.

We will report movement in either direction. A framework that only ever revises upward is not measuring anything.

Scores measure exposure to what AI can already do — not how much any particular employer has deployed. The 2023 and 2025 figures are retrospective reconstructions using the current methodology.