Pivotum

THE DEGREE RISK INDEX · FALL 2026 EDITION · FREE SAMPLER

Is Cybersecurity Safe From AI? The 2026 Degree Risk Score

The free sampler — the score by track, the six factors in brief, and the honest reasons to think twice. The complete profile goes deeper.

5.0

AI EXPOSURE · INCIDENT RESPONSE / THREAT HUNTING

where 10 is most at risk

Re-scored every six months. We publish where we might be wrong.

HIGHLIGHTING IS COLOUR-CODED

THE FINDING EXPOSURE PROTECTION METHOD & HONESTY

Quick answer: Incident response and threat hunting scores **5.0 out of 10** for AI exposure, where 10 is most at risk. Tier-one SOC analyst work scores **7.4.** Cybersecurity is widely sold as the safe technical career, and it is safer than most of tech — but "safer than software" and "safe" are different claims, and the entry tier is heavily exposed.

The short answer for parents: better than most technical paths, and not as protected as the marketing suggests. Nothing in this field scores in the low band. What protects it is an adversary who is actively trying to be unpredictable — and the roles that face that adversary directly are not the ones graduates start in.

Cybersecurity AI risk score by role

Every career in this index is scored 1–10, where **10 is most exposed** to AI. Same six factors, same weights, applied identically to a security analyst and a paramedic.

ROLE	2023	2025	FALL 2026	3-YR MOVE	BAND
Incident response / threat hunting	4.2	4.6	**5.0**	+0.8	Moderate
Security architecture	4.7	5.1	**5.5**	+0.8	Moderate
Penetration testing / offensive	4.9	5.3	**5.7**	+0.8	Moderate
GRC / compliance	5.2	5.7	**6.1**	+0.9	Moderate–High
SOC analyst (tier 1)	6.1	6.8	**7.4**	+1.3	High

2023 and 2025 figures are reconstructed using current methodology, not archived from past editions.

How that compares: the median career in this edition scores around 5.5. Entry-level software development scores **8.1.** Entry data analysis scores **8.2.** Licensed engineering scores **4.0.**

Cybersecurity is the best-scoring pure-technology track we measure — but its floor is 5.0, not 2.5. Compare licensed engineering at 4.0 or the skilled trades at 2.5.

Will AI replace cybersecurity analysts?

It is already doing most of tier-one alert triage, and struggling with everything above it.

AI IS TAKING	IT CAN'T TOUCH
Alert triage and false-positive filtering	Deciding whether this is an incident or noise
Log correlation and enrichment	Working out what an adversary is actually trying to do
Vulnerability scanning and reporting	Judgment when the attack has no precedent
Standard playbook execution	Making the call to shut down production
Compliance evidence collection	Explaining a breach to a board or a regulator
Phishing detection and classification	Adversarial thinking under time pressure

The uncomfortable symmetry: the same tools are available to attackers. Cybersecurity is the one field we score where AI improves both sides simultaneously, which is a genuinely different situation from the rest of this index.

Why is cybersecurity moderately protected? The six factors

How incident response and threat hunting rates against each. Ratings are 0–10 on each factor's own terms.

FACTOR	RATING	EFFECT ON RISK
How much of this job can AI already do?	6.0	↑ moderate exposure
How hard will it be to land that first job?	4.5	↑ moderate exposure
Does it have to be done in person, with your hands?	2.0	↓ almost none
Does someone need a human they can trust and hold responsible?	8.0	↓ strong protection
Does the law require a licensed human?	3.5	↓ weak
How often does the job hit genuinely new, high-stakes situations?	9.0	↓ strong protection

ONE FACTOR, WORKED THROUGH IN FULL

So you can see what the analysis actually looks like.

How often does the job hit genuinely new, high-stakes situations? — rated 9.0, and cybersecurity's novelty is a different kind

AI is strong on patterns and weak on genuine novelty. Most careers earn a high rating on this factor because their situations are complex — a one-off engineering site, an atypical patient, a case with no precedent.

Cybersecurity earns it because someone is deliberately manufacturing novelty against you.

That distinction matters more than it sounds. An unusual medical presentation is unusual by accident. An unusual attack is unusual *on purpose*, designed specifically to look like something benign or to exploit the gap between what a detection system was trained on and what is actually happening.

This creates a property no other profession in this index has: **the novelty is adaptive.** Whatever pattern a defensive system learns to recognize, an adversary has an incentive to stop matching it. There is no equilibrium where the patterns settle down and the work becomes routine, because the other side is paid to prevent exactly that.

That is why threat hunting rates 9.0 and holds a 5.0 score despite no physical component, no licensure, and a fully digital workflow. Almost nothing else in this index survives that combination of missing protections.

But notice where the protection does not reach. A tier-one SOC analyst is not facing adversarial novelty. They are working a queue of alerts against a playbook — pattern-matching, in a role explicitly designed to be pattern-matching, which is why it rates 4.5 rather than 9.0 and scores 7.4.

The general lesson: adversarial work resists automation better than complex work. Complexity can be learned. An opponent who changes when you learn cannot be. Any career with a genuine adversary — security, litigation, fraud investigation, negotiation, competitive strategy — carries a protection that complexity alone does not provide.

Which cybersecurity roles are safest from AI?

Ranked by exposure, safest first:

1. **Incident response / threat hunting — 5.0.** Adversarial novelty, time pressure and consequential decisions. 2. **Security architecture — 5.5.** System-level judgment where being

wrong is expensive. 3. **Penetration testing — 5.7.** Creative adversarial work, though tooling automates much of the routine scanning. 4. **GRC / compliance — 6.1.** Regulatory weight, but substantial documentation and evidence-collection content. 5. **SOC analyst tier 1 — 7.4.** Playbook-driven triage, and the traditional entry route.

Is a cybersecurity degree still worth it in 2026?

Better than most technical degrees, with one caveat that matters a great deal.

The field has genuine demand and the senior roles carry real protection. But **the standard route in — tier-one SOC work — is the most automated part of it.** That is the same entry-path problem that afflicts software, law, finance and accounting: the profession is healthy and the on-ramp is narrowing.

What still works as a route in: technical depth beyond tooling, hands-on lab and competition experience, and specialization early rather than generic security coursework. Certifications carry more weight in this field than in most, and employer-linked apprenticeship routes are more available here than in software.

Worth asking any program: what are graduates doing at year one, and does the curriculum build adversarial thinking or tool operation? **The first survives automation. The second is what is being automated.**

Common questions

Will AI replace cybersecurity jobs? It is replacing tier-one triage substantially. Incident response, threat hunting and architecture are holding, because adversarial novelty is the hardest thing for pattern-based systems.

Is cybersecurity a safe career from AI? Safer than most of technology, and not safe in absolute terms — its best track scores 5.0, against 4.0 for licensed engineering and 2.5 for the skilled trades.

Is cybersecurity safer than software engineering? Meaningfully — 5.0 for incident response against 8.1 for entry-level software development. The adversarial element is the reason.

Does AI help attackers too? Yes, and this is the field's distinctive situation. Cybersecurity is the only career we score where the same technology strengthens both sides at once, which is part of why demand is holding.

What cybersecurity jobs are safest? Incident response and threat hunting at 5.0. Anything playbook-driven scores considerably worse.

Related profiles

- **Computer science** — 8.1 entry-level, 5.4 senior. Free to read in full.
- **Data science & analytics** — 8.2 entry analyst, 6.0 ML engineering
- **Engineering** — 4.0 licensed and site-based, 6.3 desk and CAD
- **Law** — 7.9 junior associate, 4.5 trial advocacy — the other adversarial profession
- **Skilled trades** — 2.5 field and service

What's in the full cybersecurity profile

This sampler tells you where cybersecurity stands. The full profile tells you what to do about it.

	SAMPLER	FULL PROFILE
Verdict, all role scores, 3-year trend	✓	✓
Six-factor ratings	✓	✓
Reasoning behind every factor rating	one example	✓
How durable each protection is — including the adversarial dynamic		✓
The entry-path problem in detail — and the routes that still work		✓
The honest downsides — on-call burden, burnout, certification treadmill		✓
What's genuinely good about it — and who this work suits		✓
The AI-native advantage — how to prepare, and why this cohort has an opening		✓
Routes in — degree, certification, apprenticeship, competition pathways		✓
Where it leads later — architecture, leadership, consulting, research		✓
Program evaluation checklist — lab time, specialization, employer links		✓
Questions to ask a program — plus red flags		✓
Sourced further reading		✓
Discussion questions for parent and student		✓
A short version written directly to the student		✓
Technical scoring appendix		✓

GET THE FULL PROFILES

Most families are weighing two or three careers seriously, and a few more they haven't ruled out. Pick the ones you need.

1 profile	**$19**	Choose →
3 profiles	**$29**	Choose →
5 profiles	**$39**	Choose →

Each full profile includes everything in the table above, plus a short version written directly to the student and the technical scoring appendix. Spring 2027 updates of whatever you buy are included.

Read a complete profile free → We publish one profile in full, openly, so you can judge the depth before buying anything. Computer science is the one to read — it's the most surprising result in the index.

Pivotum scores thirty careers against AI exposure using a fixed, published methodology, re-scored every six months. Scores measure exposure to what AI can already do — not how much any particular employer has deployed. We publish where we might be wrong.